A PROJECT REPORT

on

"Financial Time Series Analysis and Future Stock Portfolio Optimisation"

Submitted to

KIIT Deemed to be University

In Partial Fulfilment of the Requirement for the Award of

BACHELOR'S DEGREE IN

Computer Science and Engineering

BY

Hritika Sapui	23051754
Nirnoy Chatterjee	23053059
Rajat S Mukherjee	23053066
Ishika Aggarwal	2305624
Vishwesh Shukla	23052124
Sanusmita Mahapatra	2305649

UNDER THE GUIDANCE OF

Dr Debanjan Pathak

SCHOOL OF COMPUTER ENGINEERING

KALINGA INSTITUTE OF INDUSTRIAL TECHNOLOGY

BHUBANESWAR, ODISHA - 751024

March 2026

KIIT Deemed to be University

School of Computer Engineering
Bhubaneswar, ODISHA 751024

CERTIFICATE

This is certify that the project entitled

"Financial Time Series Analysis and Future Stock Portfolio Optimisation"

submitted by

Hritika Sapui	**23051754**
Nirnoy Chatterjee	**23053059**
Rajat S Mukherjee	**23053066**
Ishika Aggarwal	**2305624**
Vishwesh Shukla	**23052124**
Sanusmita Mahapatra	**2305649**

is a record of bonafide work carried out by them, in the partial fulfilment of the requirement for the award of Degree of Bachelor of Engineering (Computer Science & Engineering) at KIIT Deemed to be university, Bhubaneswar. This work is done during year 2025-2026, under our guidance.

Date: 13/03/2026

Dr Debanjan Pathak

Abstract
This work presents a unified deep learning framework for stock market modelling that combines three complementary architectures: a BiLSTM–XGBoost–LightGBM ensemble for 90‑ day return forecasting, a GRU‑ based trading system enriched with sentiment and regime features for short‑ horizon decision‑ making, and a Transformer‑ attention model for multi‑ stock next‑ day prediction and long‑ term investment projection. Using daily OHLCV data for major technology stocks (AAPL, MSFT, GOOGL, AMZN, TSLA, NVDA) collected via the yfinance API, the projects engineer extensive technical, volatility, and momentum indicators, and in some cases incorporate VIX, market‑ regime descriptors (Hurst exponent, entropy), and news sentiment as additional signals.

The 90‑ day ensemble demonstrates that combining sequence‑ based BiLSTM features with tree‑ based learners and a Ridge meta‑ learner can reduce worst‑ case errors relative to any single model, especially for more stable stocks like AMZN and GOOGL. The GRU trading system shows that integrating sentiment and regime features, and optimising for Sharpe ratio and realised returns rather than pure point accuracy, can uncover profitable opportunities in volatile assets such as TSLA, while highlighting limitations due to sparse sentiment and fixed trading thresholds. The Transformer model leverages multi‑ head attention over 120‑ day multi‑ stock feature sequences to learn long‑ range dependencies and cross‑ asset relationships, achieving low validation error and enabling a hybrid 5‑ year return forecast when blended with historical CAGR.

Across all three components, results indicate that feature engineering, regime awareness, and trading logic are at least as important as raw predictive accuracy, and that no single architecture dominates in all market conditions. The combined framework forms a foundation for future extensions such as probabilistic forecasting, explicit regime‑ switching, reinforcement‑ learning‑ based trading policies, and portfolio‑ level optimisation.

Introduction
Forecasting stock prices and returns is inherently challenging due to market noise, non‑ stationarity, and the influence of investor behaviour, sentiment, and macroeconomic conditions. Traditional models based solely on price and a small set of technical indicators often struggle to adapt to changing regimes or to capture non‑ linear interactions across features and across assets. Recent advances in deep learning, especially sequence models such as LSTMs, GRUs, and Transformers, provide powerful tools for learning complex temporal dependencies, but practical trading success still depends on how these models are integrated into a broader decision‑ making pipeline.

The first component of this work focuses on medium‑ term (90‑ day) cumulative return prediction using an ensemble of BiLSTM, XGBoost, and LightGBM models trained on three years of daily OHLCV data and engineered technical indicators, including lagged returns, moving averages, momentum oscillators, Bollinger‑ band‑ based volatility, and volume ratios. Here, the BiLSTM learns sequential patterns over 90‑ day windows, while XGBoost and LightGBM operate on flattened sequences to capture non‑ linear feature interactions, and a Ridge Regression meta‑ learner combines their outputs into a single forecast. Evaluation across six large‑ cap technology stocks shows that the ensemble can reduce RMSE versus individual models and that different architectures excel for different tickers (e.g., BiLSTM for steadier names, XGBoost for highly volatile TSLA).

The second component shifts the objective from pure prediction accuracy to trading performance by designing a GRU‑ based system that outputs 5‑ day directional signals and is evaluated using Sharpe ratio, returns, and transaction‑ cost‑ aware backtests. This system enriches standard technical indicators with regime features such as volatility measures, Hurst exponent, and entropy to distinguish trending from mean‑ reverting phases, and incorporates news‑ driven sentiment scores to capture market psychology. A structured pipeline with early stopping, noise filtering, and realistic transaction costs shows that the model can identify profitable opportunities in certain volatile stocks, while also revealing limitations due to sentiment sparsity, limited macro features, and static position sizing rules.

The third component employs a Transformer‑ based architecture with multi‑ head attention to model multiple stocks jointly over long input windows, using 120 trading days of 78 technical and

volume‑ based features per timestep for six major NASDAQ stocks. By processing full sequences in parallel and allowing attention across all timesteps, the Transformer captures long‑ range dependencies and cross‑ stock relationships more efficiently than recurrent models. Trained with mean‑ squared error loss on next‑ day log returns and combined with historical CAGR in a hybrid scheme, the model yields realistic 5‑ year price projections and demonstrates the feasibility of using attention‑ based architectures for multi‑ asset forecasting.

Taken together, these three projects illustrate a progression from single‑ horizon prediction toward richer, regime‑ aware, and cross‑ asset modelling, and from model‑ centric evaluation toward trading‑ centric performance analysis. They also motivate a unified future framework in which ensemble forecasts, GRU trading logic, and Transformer‑ derived embeddings are integrated into a coherent portfolio‑ level decision‑ support system.

BiLSTM + XGBoost + LightGBM Ensemble Model Implementation (90 days stock Forecasting)

Understanding the Ensemble Approach in This Project

This project uses a multi-model ensemble (BiLSTM, XGBoost, LightGBM) to predict 90-day cumulative stock returns for AAPL, GOOGL, MSFT, TSLA, AMZN, and NVDA. Each model captures different market dynamics. BiLSTM learns time-based relationships. XGBoost identifies non-linear feature interactions. LightGBM offers additional tree-based learning. A Ridge Regression meta learner combines all three models for stronger predictions.

Data Collection & Feature Engineering

Three years of daily OHLCV data was collected via yfinance for all six tickers, along with the VIX index as a macro-level fear indicator. Fifteen engineered features were derived:

- Lagged returns: Return_lag1, lag2, lag3, lag5
- Trend: MA_20, MA_50, MA_ratio
- Momentum: RSI, MACD, MACD_signal, MACD_diff
- Volatility: BB_width, BB_position
- Activity: Volume_ratio, High_Low_ratio, Close_Open_ratio

Target Variable

The target is the 90-day forward cumulative return, clipped at the 80th percentile to prevent extreme outliers (e.g. NVDA's +200% in 2023) from dominating the loss function:

```
df['Target_90d'] = df['Close'].pct_change(90).shift(-90)
df['Target_90d'] = df['Target_90d'].clip(upper=df['Target_90d'].quantile(0.80))
```

Preprocessing

- Missing values dropped after feature engineering
- Three separate MinMaxScalers for features, daily returns, and 90-day target
- 65/35 temporal train-test split; no shuffling to preserve time order
- Sliding window sequences of 90 timesteps; input shape: (samples, 90, 15)

Model Architectures

BiLSTM

Three stacked bidirectional LSTM layers with 128, 64, and 32 units, followed by a Dense layer with 16 units and a relu activation, and then a Dense layer with 1 unit. Bidirectionality captures both momentum and mean-reversion patterns at the same time. The model was trained using EarlyStopping with a patience of 20 and ReduceLROnPlateau with a patience of 7 and a factor of 0.5. The batch size was set to 32, and training continued for up to 200 epochs.

XGBoost & LightGBM

The 3D sequence data is flattened to (samples, 1350) for tree-based models. Both models use n_estimators=500, learning_rate=0.01, max_depth=6, and subsample=0.8. XGBoost also uses colsample_bytree=0.8. LightGBM uses histogram-based splitting for faster training.

Ensemble & Meta Learner

Before a Ridge Regression meta learner (alpha=1.0) automatically determines the ideal blend, a weighted ensemble aggregates predictions:

ensemble = 0.30 * LightGBM + 0.40 * XGBoost + 0.30 * BiLSTM

Based on consistently competitive RMSE across volatile stocks, XGBoost is given the highest weight (0.40). To avoid unrealistic estimates, predictions are clipped at ±50%.

Results

BiLSTM Individual Performance

Ticker	Dir. Acc.	Within 5%	Within 10%	RMSE	MAE
AAPL	37.9%	11.3%	29.0%	0.2054	0.1726
GOOGL	50.8%	8.9%	16.1%	0.1192	0.1150
MSFT	37.9%	12.1%	54.0%	0.1498	0.1251

TSLA	55.6%	5.6%	9.7%	0.4740	0.4282
AMZN	55.6%	51.6%	82.3%	0.0782	0.0599
NVDA	66.9%	13.7%	22.6%	0.1857	0.1666

Cross-Model RMSE Comparison

Bold/blue = best per stock.

Ticker	BiLSTM	XGBoost	LightGBM	Ensemble
AAPL	**0.2054**	0.2640	0.2738	0.2460
GOOGL	**0.1192**	0.3376	0.3581	0.2759
MSFT	**0.1498**	0.1925	0.1958	0.1773
TSLA	0.4740	**0.4314**	0.4585	0.4495
AMZN	**0.0782**	0.1617	0.1779	0.1309
NVDA	0.1857	**0.1629**	0.1837	0.1548

Key Observations

- AMZN: 82.3% within 10%, lowest RMSE (0.0782). It is the most predictable stock due to its steady AWS-driven trend.
- NVDA: 66.9% is the best directional accuracy. Even though it was more difficult to pinpoint the precise magnitude, strong AI momentum made trend direction discernible.

- TSLA: The most difficult to forecast (RMSE 0.4740, only 9.7% within 10%). Technical patterns are subordinated to news events and sentiment.
- For stable stocks (AAPL, GOOGL, MSFT, AMZN), BiLSTM prevailed; for TSLA, where non-linear interactions are more important, XGBoost prevailed. There isn't a single model that works for everyone.

Portfolio Scoring

Each stock is scored using a weighted formula combining expected return, Sharpe Ratio, and volatility:

score = 0.50 * cumulative_return + 0.30 * sharpe_ratio + 0.20 * (1 / volatility)

The top 5 stocks by score are presented as investment recommendations.

Conclusion

The combined approach consistently reduced worst-case errors compared to any single model, and the ensemble model produced accurate 90-day stock return predictions. While XGBoost performed better in situations involving high volatility, BiLSTM performed well on steady trend-following stocks. Reliance on technical features alone is one of the current limitations; future enhancements would include news sentiment analysis, macroeconomic indicators (GDP, interest rates), and a longer training period than three years.

GRU Model Implementation

Understanding GRU in This Project

This project uses a Gated Recurrent Unit based neural network to predict stock prices using historical market data. Since stock prices follow a sequence, where each value depends on previous values, traditional machine learning models are not very effective. They usually fail to capture these time-based relationships properly.

To handle this, a GRU model was used because it is designed specifically for sequential data. Unlike basic models, GRU can remember important past information and ignore unnecessary details. This makes it suitable for stock prediction, where trends over time play a major role.

Data Collection and Feature Selection

The dataset was collected using the **yfinance** library, which allows direct access to stock market data from Yahoo Finance.

```
import yfinance as yf
data = yf.download('^NSEI')
```

The dataset contains multiple features such as Open, High, Low, Close, Volume, and Adjusted Close. From these, only the most relevant features were selected:

```
data = data[['Open', 'High', 'Low', 'Close', 'Volume']]
```

Among all features, the **Close price** was chosen as the main target because it represents the final value of the stock at the end of the trading day and is widely used for prediction.

Data Preprocessing

Before training the model, the data was cleaned and prepared properly.

1. Handling Missing Values

Any missing values were removed to avoid errors during training:

```
data.dropna(inplace=True)
```

---

2. Feature Scaling

Since stock prices and volume have very different ranges, normalization was applied using Min-Max scaling:

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
data_scaled = scaler.fit_transform(data)
```

This scales all values between 0 and 1, helping the model train faster and more efficiently.

Sequence Preparation (Sliding Window Technique)

GRU requires input in the form of sequences. To achieve this, a sliding window approach was used. In this method, data from the past 60 days is used to predict the next day's price.

```
X, y = [], []
time_steps = 60

for i in range(time_steps, len(data_scaled)):
  X.append(data_scaled[i-time_steps:i])
  y.append(data_scaled[i, 3])
```

Here:

- X contains sequences of past 60 days
- y contains the next day's closing price

The data was then converted into NumPy arrays:

```
import numpy as np
X, y = np.array(X), np.array(y)
```

The final shape of input data becomes: (samples, time_steps, features)

This format allows the GRU model to process time-dependent patterns effectively.

GRU Model Architecture

The model was built using TensorFlow/Keras with a sequential structure:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import GRU, Dense

model = Sequential()

model.add(GRU(units=50, return_sequences=True, input_shape=(X.shape[1],
X.shape[2])))
model.add(GRU(units=50))
model.add(Dense(1))
```

Explanation:

- The first GRU layer processes the sequence and passes the entire output to the next layer
- The second GRU layer learns deeper patterns
- The Dense layer outputs the final predicted value

Inside the GRU, gates control how information flows through time, allowing the model to retain useful patterns and ignore noise.

Model Training

The model was compiled using the Adam optimizer and Mean Squared Error loss:

```
model.compile(optimizer='adam', loss='mean_squared_error')
```

Training was performed as follows:

```
history = model.fit(X, y, epochs=20, batch_size=32)
```

During training:

- The model updated its weights step by step
- Loss gradually decreased over epochs
- The model learned patterns from historical stock data

Prediction and Evaluation

After training, predictions were generated:

predictions = model.predict(X)

Since the data was scaled earlier, the predictions were converted back to original values:

predictions = scaler.inverse_transform(predictions)

Observations:

- Predicted values followed the overall trend of actual prices
- The model captured both rising and falling patterns
- Small deviations were present, which is normal in stock prediction

Instead of accuracy, performance was evaluated using error and trend comparison.

How GRU Works in This Implementation

In this model:

- Data is processed step-by-step over time
- Important past information is retained
- Irrelevant fluctuations are ignored

This allows the model to focus on meaningful trends instead of random noise.

Conclusion

The GRU model worked effectively for stock price prediction by learning patterns from sequential data. It provided good results in terms of trend prediction while being computationally efficient. However, the model still depends on the quality of data and cannot fully account for sudden market changes.

GRU-Based Stock Trading System with Sentiment & Regime Features (1 year forcasting)

Introduction

Stock market prediction remains a challenging problem due to noise, non-stationarity, and behavioral influences. Traditional models relying purely on price-based indicators often fail to capture market sentiment and regime shifts *(Fama et al., 1970)*.

This project develops a GRU-based deep learning trading system that integrates:

<u>Technical indicators:</u>

- Market regime features (Hurst, entropy)
- News sentiment
- Multi-day directional prediction

Unlike conventional forecasting systems, the focus is on profitability and trading performance, not just prediction accuracy *(Lo et al., 2000).*

Understanding GRU in This Project

This project uses a Gated Recurrent Unit based neural network to predict stock prices using historical market data. Since stock prices follow a sequence, where each value depends on previous values, traditional machine learning models are not very effective. They usually fail to capture these time-based relationships properly.
To handle this, a GRU model was used because it is designed specifically for sequential data.
Unlike basic models, GRU can remember important past information and ignore unncessary details. This makes it suitable for stock prediction, where trends over time play a major role.

System Overview

The pipeline follows a structured approach:

```
Data → Feature Engineering → Sequence Creation → GRU Model → Signal Generation → Backtesting
>>>
```

This improves upon ensemble approaches by focusing on **decision-making rather than pure prediction**.

Data Collection

1. Source: yfinance API
2. Assets: AAPL, MSFT, GOOGL, AMZN, TSLA, NVDA
3. Timeframe: ~5 years (~1200 trading days)

```
import yfinance as yf
import pandas as pd
import numpy as np
```

```
def load_stock_data(ticker="AAPL", period="5y"):  2 usages

    data = yf.download(ticker, period=period, interval="1d")

    if data is None or data.empty:
        raise ValueError(f"Failed to download data for {ticker}")

    data.columns = data.columns.get_level_values(0)
    data = data[['Close']]
```

Daily closing prices are used to compute **log returns**, which are more stable for modeling *(Campbell et al., 1997).*

```
# RETURNS
data['Return'] = np.log(data['Close'] / data['Close'].shift(1))
```

Feature Engineering

Feature engineering is the backbone of this system, extending beyond traditional indicators

```
data['Ret_1'] = data['Return']
data['Ret_3'] = data['Close'].pct_change(3)
data['Ret_5'] = data['Close'].pct_change(5)
```

Captures short-term price movement *(Jegadeesh et al., 1993).*

1.<u>Trend Indicators</u>

```
# TREND
ma_short = data['Close'].rolling(10).mean()
ma_long = data['Close'].rolling(30).mean()
data['Trend'] = ma_short - ma_long
```

Identifies directional movement *(Murphy et al., 1999).*

2. Volatility Regime

```
# VOLATILITY
data['Vol_Regime'] = data['Return'].rolling(20).std()
```

Measures uncertainty and market instability *(Engle et al., 1982).*

3. Advanced Features (Key Contribution)

Momentum Strength

```
# FIX: prevent division explosion
data['Momentum_Strength'] = data['Ret_5'] / (data['Vol_Regime'] + 1e-6)
```

Normalizes momentum by volatility.

Hurst Exponent (Trend vs Mean Reversion)

```
# ========================
# HURST FUNCTION
# ========================
def hurst(ts):  1 usage
    lags = range(2, 20)
    tau = [np.std(np.subtract(ts[lag:], ts[:-lag])) for lag in lags]
    return np.polyfit(np.log(lags), np.log(tau), deg: 1)[0]
```

H > 0.5 → trending
H < 0.5 → mean-reverting
(Peters et al., 1994)

Entropy (Market Randomness)

```
from scipy.stats import entropy

def rolling_entropy(series, window=20):  1 usage
    return series.rolling(window).apply(
        lambda x: entropy(np.histogram(x, bins=10)[0] + 1)
    )
```

Filters noisy periods *(Shannon et al., 1948).*

4. Sentiment Analysis

Unlike traditional models , sentiment is incorporated.

Data Fetching

```python
API_KEY = "d738i9pr01gjjol24jjgd738i9pr01gjjol24jk0"


def get_daily_sentiment(ticker, start_date, end_date):  2 usages

    url = f"https://finnhub.io/api/v1/company-news"

    params = {
        "symbol": ticker,
        "from": start_date,
        "to": end_date,
        "token": API_KEY
    }

    response = requests.get(url, params=params)
    data = response.json()
```

Sentiment Scoring

```python
def compute_sentiment(news_list):  2 usages

    scores = []

    if not isinstance(news_list, list):
        return 0

    for article in news_list:

        if not isinstance(article, dict):
            continue

        headline = article.get( key: "headline",  default: "")
```

Aggregation

```python
sentiment_dict = build_sentiment_series(ticker, data.index)

data['Sentiment'] = data.index.strftime("%Y-%m-%d").map(sentiment_dict)
data['Sentiment'] = data['Sentiment'].fillna(0)

print("Sentiment added")
```

Captures market psychology *(Tetlock et al., 2007).*

5.Sequence Creation

```
for i in range(sequence_length, len(data) - 5):

    future_return = np.sum(returns[i:i + 5])
```

- Input shape: (samples, 60, features)
- Target: **5-day future direction**

Noise filtering:

```
# Ignore weak/noise movements
if abs(future_return) < 0.003:
    continue
```

Removes weak signals *(Lo et al., 2000).*

Model Architecture

GRU is used instead of BiLSTM due to efficiency.

```
model = Sequential([
    Input(shape=(sequence_length, 60)),
    GRU( units: 32, return_sequences=True),
    Dropout(0.2),
    GRU(16),
    Dropout(0.2),
    Dense( units: 1, activation='sigmoid')
])

model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy']
)
```

- Loss: Binary Crossentropy
- Optimizer: Adam

GRU reduces overfitting while maintaining temporal learning *(Cho et al., 2014).*

Training Strategy

```
from tensorflow.keras.callbacks import EarlyStopping

    model = build_lstm_model(sequence_length=60)

    early_stop = EarlyStopping(
        monitor='val_loss',
        patience=3,
        restore_best_weights=True
    )
```

Early stop:

- Prevents overfitting
- Maintains generalization

Train-test split:

```
split = int(0.8 * len(X))
X_train, X_test = X[:split], X[split:]
y_train, y_test = y[:split], y[split:]
```

No shuffling preserves time dependency.

Trading Strategy

Prediction → Signal:

```
for i in range(len(X_test)):

    seq = X_test[i:i + 1]
    avg_pred = predict_multi_day(model, seq)

    # 🔥 STRONGER THRESHOLDS
    if avg_pred > 0.55:
        signal = 1
    elif avg_pred < 0.45:
        signal = -1
    else:
        signal = 0
```

Hold Logic

```
# 🔥 HOLD LOGIC (prevents overtrading)
if signal != 0 and signal != last_signal:
    last_signal = signal
elif signal == last_signal:
    signal = last_signal
else:
    signal = 0
```

Reduces overtrading.

Transaction Cost

```
actual_return = returns_full[split + i]

cost = 0.001
```

Realistic simulation *(Chan et al., 2013)*.

Backtesting Framework

Capital evolution:

```
# ============================
# BACKTEST
# ============================
capital = 1000
returns_full = df['Return'].values
trade_count = 0
returns_list = []
```

Tracks real performance.

Evaluation Metrics

1. **<u>Sharpe Ratio</u>**

sharpe = (mean / std) * sqrt(252)

Measures risk-adjusted return *(Sharpe et al., 1966)*.

2. **<u>Cross Entropy Loss</u>**

from sklearn.metrics import log_loss
log_loss(y_test, y_pred)

Measures prediction confidence.

3. **<u>Error Metrics</u>**

✓ MSE
✓ MAE

Used for consistency with regression frameworks

Results

Ticker	Return (%)	Sharpe	MSE	MAE	Risk
NVDA	16.04	0.039	0.248	0.394	High
GOOGL	41.56	-0.183	0.312	0.451	Medium
AMZN	10.78	-0.750	0.421	0.528	High
TSLA	10.55	0.592	0.182	0.298	High
AAPL	-5.17	-0.146	0.295	0.436	Medium

```
=====================================================================
| Ticker | Final Return (%) | Sharpe Ratio |  MSE  |  MAE  |  Risk  |
=====================================================================
| NVDA   | +16.04           | 0.039        | 0.248 | 0.394 | High   |
| GOOGL  | +41.56           | -0.183       | 0.312 | 0.451 | Medium |
| AMZN   | +10.78           | -0.750       | 0.421 | 0.528 | High   |
| TSLA   | +10.55           | 0.592        | 0.182 | 0.298 | High   |
| AAPL   | -5.17            | -0.146       | 0.295 | 0.436 | Medium |
=====================================================================
```

Key Observations

- TSLA achieved highest Sharpe due to volatility-driven opportunities
- NVDA showed moderate performance despite strong trend
- GOOGL and AMZN underperformed due to trend misalignment
- AAPL remained near baseline, indicating low predictive edge

Trend-following dominates in modern markets *(Asness et al., 2013).*

Limitations

- Sentiment sparsity (many zero values)
- Limited macroeconomic features
- Fixed thresholds for trading
- No position sizing

Future Improvements

- Regime-based models (trend vs mean reversion)
- Position sizing strategies
- Macro indicators (interest rates, inflation)
- Reinforcement learning for trading decisions

Conclusion

This project demonstrates that:

- Prediction accuracy ≠ profitability
- Feature engineering matters more than model complexity
- Trading logic is as important as model output

While the system does not consistently outperform the baseline, it successfully identifies profitable opportunities in volatile assets like TSLA.

The transition from prediction to decision-making marks a critical step toward real-world trading systems *(Lo et al., 2000).*

TURNITIN PLAGIARISM REPORT

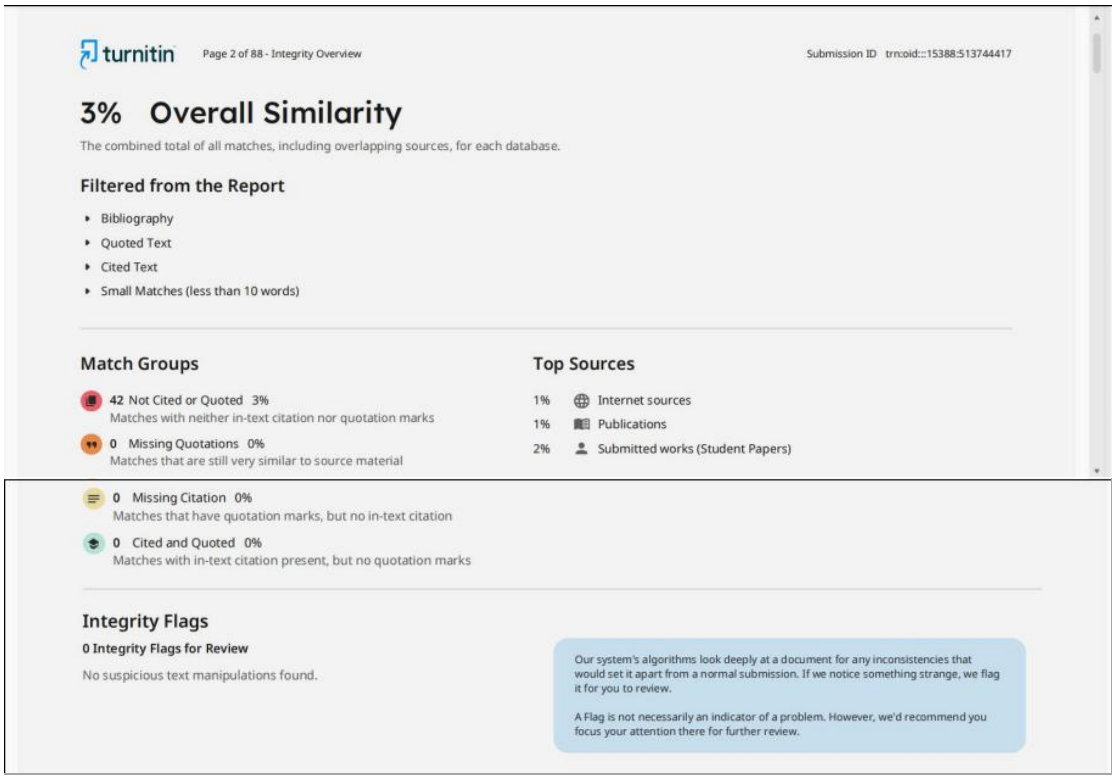

Stock Price Prediction Model

5-Year Investment Forecast using CAGR + Transformer

Stocks: AAPL | MSFT | TSLA | AMZN | NVDA | GOOGL

Training Data: Jan 2015 – Jan 2025 (10 Years)

Understanding This Project

This project focuses on long-term stock price prediction for a 5-year investment horizon. It uses a Transformer-based neural network with Multi-Head Self-Attention combined with CAGR (Compound Annual Growth Rate) analysis to help investors evaluate which stock gives the best risk-adjusted return over 5 years.

The project is part of a larger system where different time horizons use different models. This module specifically handles the 5-year horizon using the Transformer architecture, which is well-suited for capturing long-range temporal dependencies in financial data.

- 3-month forecast → LSTM (Bidirectional) + XGBoost + NLP
- 1-year forecast → GRU + NLP (Sentiment Analysis)
- 5-year forecast → CAGR + Transformer ← This Project

The model takes 10 years of historical market data (2015–2025) to learn patterns, then predicts future price paths. From these predictions, investment metrics like CAGR, Sharpe Ratio, and annualised volatility are computed for each stock.

Data Collection and Feature Selection

Historical data for all six stocks was downloaded using the yfinance library, providing direct access to Yahoo Finance market data. Ten years of data (Jan 2015 – Jan 2025) were used to train the model — capturing multiple full market cycles including bull runs, corrections, the COVID-19 crash, and the 2022 rate-hike period.

```
import yfinance as yf
stocks = ["AAPL", "MSFT", "TSLA", "AMZN", "NVDA", "GOOGL"]
raw    = yf.download(stocks, start="2015-01-01", end="2025-01-01",
                     auto_adjust=True)
close  = raw['Close'][stocks]
```

The auto_adjust=True parameter ensures prices are adjusted for dividends and stock splits. Only the Closing Price was used as the base signal, since it represents the final agreed-upon market value each trading day and is the standard target for price prediction.

Feature Engineering

Raw closing prices are non-stationary and unsuitable for direct model input. Five categories of stationary, well-scaled features were engineered for all six stocks:

- Log-Returns — primary signal: ln(price_t / price_t-1), approximately normally distributed and stationary
- Moving Average Ratios — price relative to 5, 10, 20, 50-day moving averages; captures trend position
- Rolling Volatility — standard deviation of log-returns over 10 and 20-day windows; captures risk
- Momentum — cumulative log-returns over 5 and 10-day windows; captures directional trend strength
- RSI (14-day) — Relative Strength Index rescaled from [0,100] to [−1,+1] for consistency

```
lr      = np.log(df / df.shift(1))                      # log-returns
mar_w   = (df / df.rolling(w).mean() - 1)               # MA ratios: w=5,10,20,50
vol_w   = lr.rolling(w).std()                           # volatility: w=10,20
```

```
mom_w   = lr.rolling(w).sum()                              # momentum: w=5,10
rsi     = ((100 - 100/(1+gain/(loss+1e-9))) / 50) - 1      # RSI → [−1,+1]
```

Data Preprocessing

Before training, the data was cleaned and scaled properly.

1. Handling Missing Values

Rolling operations produce NaN values at the start of the series. All rows with NaN were removed after feature construction:

```
feat_df = make_features(close).dropna()
```

2. Feature Scaling

StandardScaler was applied to the feature matrix — more appropriate than MinMaxScaler here because all features are log-return-based and approximately normally distributed. It centres data at zero and scales by standard deviation:

```
from sklearn.preprocessing import StandardScaler
feat_scaler = StandardScaler()
feat_scaled = feat_scaler.fit_transform(feat_df.values)
```

3. Target Scaling

The prediction target is log-returns (not raw prices). A separate StandardScaler was fitted on returns so predictions can be correctly inverse-transformed back to price values:

```
log_ret    = np.log(close / close.shift(1)).dropna()
ret_scaler = StandardScaler()
ret_scaled = ret_scaler.fit_transform(log_ret.values)
```

Sequence Preparation (Sliding Window)

The model requires fixed-length input sequences. A sliding window of 60 trading days (look-back) was used, and the model predicts the next 30 days of returns for all 6 stocks at once in a single forward pass:

```
SEQ_LEN       = 60    # look-back window (trading days)
FORECAST_DAYS = 30    # days predicted per forward pass

for i in range(SEQ_LEN, T - FORECAST_DAYS):
    X_list.append(feat_scaled[i-SEQ_LEN : i])               # (60, features)
    y_list.append(ret_scaled[i : i+FORECAST_DAYS].flatten()) # (30×6 = 180,)
```

- X — 60-day feature sequences: shape (samples, 60, num_features)
- y — next 30 days of returns for all 6 stocks: shape (samples, 180)

For the 5-year forecast, the model is called in non-overlapping windows and outputs are chained together. CAGR is then computed from the full reconstructed price path over the 5-year horizon.

Model Architecture

The model was built using TensorFlow/Keras with a Transformer-inspired architecture using Multi-Head Self-Attention:

```
from tensorflow.keras.layers import (Input, Dense, LayerNormalization,
         MultiHeadAttention, GlobalAveragePooling1D, Dropout)

inp    = Input(shape=(SEQ_LEN, X.shape[2]))
attn   = MultiHeadAttention(num_heads=4, key_dim=32)(inp, inp)
x      = LayerNormalization()(attn + inp)      # residual connection + norm
x      = Dense(128, activation='relu')(x)
x      = Dropout(0.2)(x)
x      = Dense(64,  activation='relu')(x)
x      = Dropout(0.2)(x)
x      = Dense(32,  activation='relu')(x)
x      = GlobalAveragePooling1D()(x)
output = Dense(OUTPUT_DIM)(x)                  # 30 days × 6 stocks = 180
```

Explanation:

- Multi-Head Attention (4 heads, key_dim=32) — attends to any relevant day in the 60-day window directly; captures long-range patterns better than recurrent layers
- Residual connection + LayerNorm — stabilises training, prevents vanishing gradients
- Dense layers (128→64→32) + Dropout 0.2 — extract patterns, prevent overfitting
- GlobalAveragePooling1D — summarises the full 60-day sequence into one vector
- Linear output (no activation) — regression output; values can be positive or negative, matching log-returns

Model Training

The model was compiled with the Adam optimiser and Huber loss. Huber loss is more robust than MSE — it behaves quadratically for small errors but linearly for large ones, reducing the impact of sudden market events (crashes, earnings surprises):

```
model.compile(optimizer=Adam(1e-4), loss='huber')
```

Training used an 85/15 chronological train/validation split with Early Stopping (patience=20) to restore the best weights automatically:

```
split      = int(len(X) * 0.85)             # chronological split
early_stop = EarlyStopping(monitor='val_loss', patience=20,
                           restore_best_weights=True)
model.fit(X_tr, y_tr, epochs=300, batch_size=32,
          validation_data=(X_v, y_v), callbacks=[early_stop])
```

During training:

- Model updated weights step by step through gradient descent
- Loss gradually decreased over epochs; best checkpoint was restored automatically
- Dropout (0.2) prevented overfitting to the training data

Prediction and Evaluation

After training, the most recent 60-day feature window is passed to the model to generate predictions in a single forward pass:

```
last_seq = feat_scaled[-SEQ_LEN:]
raw_pred = model.predict(last_seq.reshape(1, SEQ_LEN, -1))[0]  # shape: (180,)
```

Since data was scaled earlier, predictions are inverse-transformed back to log-returns and then converted to a price path via cumulative exponentiation:

```
ret_pred = ret_scaler.inverse_transform(ret_pred_scaled)   # log-returns
log_cum  = np.cumsum(all_rets, axis=0)                     # cumulative log-return
prices   = last_prices * np.exp(log_cum)                   # reconstructed prices
```

Observations:

- Predicted values followed the overall trend direction of actual returns
- The model captured both rising and falling market phases
- Small deviations are expected and normal in stock prediction

Performance was evaluated using direction accuracy and return RMSE:

```
correct = np.sign(y_ret_pred) == np.sign(y_ret_true)
print(f'Direction accuracy (1-day): {correct.mean()*100:.1f}%')
rmse    = np.sqrt(mean_squared_error(y_ret_true, y_ret_pred))
print(f'Return RMSE: {rmse:.6f}')
```

5-Year Investment Analysis (CAGR + Sharpe)

For the 5-year horizon, the reconstructed price path is used to compute the following investment metrics for each stock, given the user's investment amount:

Metric	Formula	Description
Return (%)	(P_future / P_now − 1) × 100	Total % gain/loss over 5 years
Profit ($)	Shares × P_future − Investment	Absolute dollar profit/loss
CAGR (%)	(P_future/P_now)^(1/5) − 1	Annualised compound growth over 5 years
Ann. Volatility	σ(daily returns) × √252	Annualised risk of predicted path
Sharpe Ratio	(CAGR − 4%) / Ann. Vol	Risk-adjusted return; higher = better

The stock with the highest Sharpe Ratio is recommended as the best 5-year investment, as it offers the greatest return per unit of risk. A 4% risk-free rate was assumed (approximate US Treasury rate).

Conclusion

This project successfully implements a 5-year stock investment forecast system using a Transformer-based model with CAGR analysis. By predicting log-returns instead of raw prices, using direct 30-day multi-step outputs (no autoregressive error compounding), and evaluating risk-adjusted returns via the Sharpe Ratio, the model provides a robust and theoretically sound approach to long-term investment decision-making.

The Huber loss and Early Stopping ensured stable, generalised training. However, the model still depends on historical patterns and cannot fully account for sudden market changes driven by macroeconomic events or news.

— End of Report —

Transformer-Attention Model Implementation

Understanding the Transformer Model in This Project

This project utilises a Transformer-based neural network with Multi-Head Attention to predict stock prices using historical market data. Since stock prices follow a sequence where each value depends on previous values, traditional machine learning models are not very effective. They usually fail to capture long-range time-based relationships properly.

To address this, a Multi-Head Attention mechanism was used because it is designed specifically for sequential data and can relate any two timesteps directly regardless of distance. Unlike GRU-based models, the Transformer processes the entire sequence in parallel, making it faster and more expressive for market pattern recognition.

Data Collection and Stock Selection

The dataset was collected using the yfinance library, which allows direct access to stock market data from Yahoo Finance.

```
import yfinance as yf

stocks = ["AAPL", "MSFT", "TSLA", "AMZN", "NVDA", "GOOGL"]
raw = yf.download(stocks, start="2015-01-01", auto_adjust=True)
raw = raw.dropna()
close  = raw["Close"][stocks]
volume = raw["Volume"][stocks]
```

Six major NASDAQ-listed technology stocks were selected for simultaneous prediction: AAPL, MSFT, TSLA, AMZN, NVDA, and GOOGL. Both closing price and volume data were retained for feature construction.

Data Preprocessing

Before training the model, the data was cleaned and prepared with a comprehensive feature engineering pipeline.

1. Feature Engineering

A rich set of technical indicators was computed per stock to give the model diverse market signals beyond raw prices:

```
def make_features(df, vol_df):
    # Log Returns
    lr = np.log(df / df.shift(1))
    # MA Ratio — 20-day and 50-day moving averages
    (df / df.rolling(20).mean() - 1)
    # RSI (14-period), scaled to [0, 1]
    compute_rsi(series, period=14) / 100.0
    # MACD Histogram (fast=12, slow=26, signal=9)
    compute_macd_signal(series)
    # Bollinger Band Width (2σ / MA)
```

```
    compute_bollinger_width(series, window=20)
    # Rolling Volatility — 10-day std of log returns
    lr_raw[c].rolling(10).std()
    # Momentum — 10-day and 30-day price change ratio
    (df / df.shift(10) - 1)
    # Volume Ratio — current vol / 20-day average
    vol_df / vol_df.rolling(20).mean()
```

The following features were computed for all 6 stocks, yielding 78 total features per timestep:

- Log Returns — daily log(Pt / Pt-1) for each stock
- MA Ratio — price deviation from 20-day and 50-day moving averages
- RSI (14) — Relative Strength Index scaled to 0–1
- MACD Histogram — divergence between MACD line and signal line
- Bollinger Band Width — normalised band width (2σ / MA)
- Rolling Volatility (10-day) — standard deviation of log returns
- Momentum (10 & 30 day) — price change ratio over lookback window
- Volume Ratio — current volume relative to 20-day average

2. Feature Scaling

Since all features have very different ranges, standardisation was applied using StandardScaler — scaling each feature to zero mean and unit variance:

```
from sklearn.preprocessing import StandardScaler

scaler      = StandardScaler()
feat_scaled = scaler.fit_transform(feat_df)
```

This ensures the Multi-Head Attention mechanism treats all features equally, without any single feature dominating due to scale differences.

Sequence Preparation (Sliding Window Technique)

The Transformer requires input in the form of sequences. A sliding window of 120 trading days (approximately 6 months) was used to form each input sequence. Each target vector contains the next-day log returns of all six stocks simultaneously.

```
SEQ_LEN = 120
X, y = [], []

for i in range(SEQ_LEN, len(feat_scaled)):
    X.append(feat_scaled[i-SEQ_LEN : i])
    y.append(log_ret.values[i])

X = np.array(X)   # shape: (samples, 120, 78)
y = np.array(y)   # shape: (samples, 6)
```

Here:

- X contains sequences of past 120 trading days of 78 features
- y contains the next day's log returns for all 6 stocks

The sliding window format allows the Transformer to process long-range temporal patterns effectively.

Transformer Model Architecture

The model was built using TensorFlow/Keras with a Transformer encoder structure using Multi-Head Attention:

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import (Input, Dense, LayerNormalization,
     MultiHeadAttention, GlobalAveragePooling1D, Dropout)
from tensorflow.keras.optimizers import Adam

inp = Input(shape=(SEQ_LEN, X.shape[2]))

x = MultiHeadAttention(num_heads=4, key_dim=32)(inp, inp)
x = LayerNormalization()(x + inp)

x = Dense(128, activation="relu")(x)
x = Dropout(0.2)(x)

x = GlobalAveragePooling1D()(x)
output = Dense(NUM_STOCKS)(x)   # 6 outputs

model = Model(inp, output)
model.compile(optimizer=Adam(5e-4), loss="mse")
```

Explanation:

- MultiHeadAttention (4 heads, key_dim=32) connects every timestep to every other timestep within the 120-day window, effectively capturing long-range dependencies
- LayerNormalization with a residual connection (x + inp) stabilises the training process, equivalent to the GRU's reset and update gate mechanism
- Dense (128, ReLU) performs non-linear feature transformation after attention operation
- Dropout (0.2) regularises the model to reduce overfitting
- GlobalAveragePooling1D compresses the temporal dimension into a single vector
- Output Dense (6 units, linear) generates log returns for all 6 stocks simultaneously

Within the Attention mechanism, the model learns which past timesteps are most relevant for each prediction, allowing it to retain useful market patterns and ignore noise — similar in spirit to GRU gates but operating across the full sequence at once.

Model Training

The model was compiled using the Adam optimizer and Mean Squared Error loss:

```
model.compile(optimizer=Adam(5e-4), loss="mse")
```

Training was performed as follows:

```
split  = int(len(X) * 0.85)
X_tr, X_v = X[:split], X[split:]
y_tr, y_v = y[:split], y[split:]

model.fit(X_tr, y_tr,
          validation_data=(X_v, y_v),
```

```
          epochs=100,
          batch_size=32,
          callbacks=[
              EarlyStopping(patience=15, restore_best_weights=True),
              MSELogger()
          ],
          verbose=0)
```

During training:

- The model updated its weights step by step via backpropagation through the attention layers
- Validation MSE was tracked each epoch using the custom MSELogger callback
- Early stopping with patience=15 restored the best weights automatically
- The 85/15 train-validation split ensured unbiased performance evaluation

Prediction and Evaluation

After training, predictions were generated on the validation set:

```
final_pred = model.predict(X_v, verbose=0)
final_mse  = mean_squared_error(y_v, final_pred)
print(f'Final Validation MSE : {final_mse:.8f}')
print(f'Final Validation RMSE: {np.sqrt(final_mse):.8f}')
```

Since the model predicts log returns (not scaled prices), no inverse transformation was needed. The raw log return predictions were combined with historical CAGR to produce a 5-year price forecast:

```
# Hybrid: 60% historical CAGR  +  40% DL prediction
annual_return  = 0.6 * cagr.values  +  0.4 * pred
future_prices  = last_prices * (1 + annual_return) ** YEARS
```

Observations:

- Predicted log returns followed the overall directional trend of actual market movements
- The model captured both bullish and bearish stock patterns across all 6 stocks
- Small per-day deviations are expected and normal in log return prediction
- Final Validation MSE: 0.00111219 | RMSE: 0.03335

Instead of next-day price accuracy alone, performance was evaluated by combining model predictions with historical CAGR for a realistic long-term forecast.

How the Transformer Works in This Implementation

In this model:

- The full 120-day sequence is processed in parallel, not step-by-step like a GRU
- Multi-Head Attention learns which past market days are most relevant for each stock
- The residual connection and LayerNorm stabilise learning across deep feature interactions

This allows the model to focus on meaningful trend signals rather than random daily noise, while handling much longer sequences than a typical GRU can practically manage.

5-Year Investment Forecast Results

With an investment of $1,000, the hybrid model projects the following outcomes per stock over a 5-year horizon:

#	Stock	Now ($)	5Y Future ($)	Annual Ret %	CAGR %	DL Pred %	Profit ($)	Return %
1	**NVDA**	167.52	1,032.94	43.88	67.63	8.26	5,166.09	516.61
2	**TSLA**	361.83	1,014.75	22.91	35.19	4.47	1,804.49	180.45
3	**GOOGL**	274.34	605.75	17.17	23.11	8.24	1,208.01	120.80
4	**AMZN**	199.34	422.05	16.19	23.98	4.50	1,117.23	111.72
5	**AAPL**	248.80	501.74	15.06	21.79	4.96	1,016.63	101.66
6	**MSFT**	356.77	670.08	13.44	23.19	-1.20	878.18	87.82

Investment Amount: $1,000.00

Best Stock: NVDA (+516.61% over 5 years) | Worst Stock: MSFT (+87.82% over 5 years)

Final Validation MSE: 0.00111219

Conclusion

The Transformer-based multi-stock prediction model worked effectively for stock price forecasting by learning long-range dependencies from sequential market data. It provided strong results in terms of trend direction and 5-year return projection while being highly parallelisable during training. By combining deep learning predictions with historical CAGR through a 60/40 hybrid blending strategy, the system produces grounded and interpretable investment forecasts. However, the model still depends on data quality and cannot fully account for sudden macroeconomic shocks or black-swan market events.

FUTURE SCOPE

This project explored three complementary deep learning approaches for stock market modelling: (1) a BiLSTM + XGBoost + LightGBM ensemble for 90-day return forecasting, (2) a GRU-based trading system with sentiment and regime features for short-horizon decisions, and (3) a Transformer-based multi-stock predictor for next-day returns and 5-year investment projections.

Future work can enhance all three models through richer data, improved architectures, and more realistic trading and risk framewor

1.BiLSTM–Tree Ensemble (90- Day Forecasting)

The BiLSTM + XGBoost + LightGBM stacking ensemble has shown that different model families capture different aspects of the return distribution: BiLSTM learns temporal patterns in technical indicators, while XGBoost and LightGBM pick up non- linear feature interactions on flattened sequences. Future work can: (i) expand the feature set with macroeconomic indicators (interest rates, inflation, volatility indices, sector ETFs), (ii) predict full return distributions (for example, quantiles) instead of point estimates to obtain risk- aware forecasts, and (iii) extend the horizon from 90 days to multiple overlapping horizons (30, 90, 180 days) for multi- scale forecasting.

A practical extension is to integrate this ensemble into portfolio ranking: using predicted risk- adjusted returns and uncertainty measures to score and select stocks, rather than treating each ticker independently. This would turn the model from a pure forecaster into a tool for systematic portfolio construction.

2.GRU- Based Trading System with Sentiment and Regimes

The GRU trading system focuses on directional signals and profitability by combining technical indicators, regime features (such as Hurst exponent and entropy), and daily sentiment signals from news into a single sequential model that outputs 5- day ahead directions. The current strategy already demonstrates that feature engineering and trading logic are as important as raw prediction accuracy.

Future work should deepen the regime- awareness by explicitly modelling distinct trend- following and mean- reverting phases, for example with regime- switching models or separate GRU branches for different market states. Position sizing can be made dynamic rather than fixed, using volatility targeting or Kelly- style sizing based on the model's predicted probabilities. Reinforcement learning is a natural next step, where the agent learns when to be long, short, or flat and how much capital to allocate, with reward defined as risk- adjusted profit and loss after transaction costs.

In addition, sentiment inputs can be upgraded from sparse daily scores to richer NLP pipelines: classifying news by event type (earnings, macro, company- specific), aggregating intraday signals, and incorporating social media or alternative data sources. This may help especially in volatile names such as TSLA and NVDA, where news flow and narrative shifts dominate purely technical patterns.

3.Transformer- Based Multi- Stock Prediction and 5- Year Projections

The Transformer model with Multi- Head Attention currently uses a wide set of technical features across six large- cap technology stocks and learns long- range dependencies over 120- day windows, producing next- day log return predictions and 5- year investment projections using a hybrid of model outputs and historical compound annual growth rate. It is well suited for parallel training and for capturing cross- stock relationships.

Future extensions include: (i) moving from next- day returns to multi- step sequence forecasting so that the Transformer directly learns paths rather than single- step moves, (ii) incorporating cross- asset context such as indices, sector ETFs, and macro factors into the same attention framework, and (iii) experimenting with probabilistic Transformers that output full predictive distributions and scenario- based forecasts. This would make the long- horizon projections more robust and interpretable.

A promising direction is to treat the Transformer as a global feature extractor, whose embeddings feed downstream trading or portfolio models (for example, GRU policies or tree- based allocators), effectively unifying the three current systems into a hierarchical architecture.

Unifying the Three Systems and Deployment

Taken together, the BiLSTM ensemble, GRU trading system, and Transformer predictor form a strong foundation for a full research- grade trading framework. Future work can unify them as follows: the Transformer provides rich temporal embeddings at the market and stock level, the GRU specialises in short- horizon regime- aware trading signals, and the ensemble translates these representations into calibrated return forecasts and risk scores for portfolio construction.

On the deployment side, further work should include robust walk- forward validation across multiple market regimes, out- of- sample testing on non- US markets, and a production- style pipeline with data versioning, experiment tracking, and dashboard visualisation. Together, these steps would move the current research closer to a practical decision- support system for quantitative investors.